

Practice Lab: Neural Networks for Handwritten Digit Recognition, Multiclass

In this exercise, you will use a neural network to recognize the hand-written digits 0-9.

Outline

- [1 - Packages](#)
- [2 - ReLU Activation](#)
- [3 - Softmax Function](#)
 - [Exercise 1](#)
- [4 - Neural Networks](#)
 - [4.1 Problem Statement](#)
 - [4.2 Dataset](#)
 - [4.3 Model representation](#)
 - [4.4 Tensorflow Model Implementation](#)
 - [4.5 Softmax placement](#)
 - [Exercise 2](#)

NOTE: To prevent errors from the autograder, you are not allowed to edit or delete non-graded cells in this notebook. Please also refrain from adding any new cells. Once you have passed this assignment and want to experiment with any of the non-graded code, you may follow the instructions at the bottom of this notebook.

1 - Packages

First, let's run the cell below to import all the packages that you will need during this assignment.

- [numpy](#) is the fundamental package for scientific computing with Python.
- [matplotlib](#) is a popular library to plot graphs in Python.
- [tensorflow](#) a popular platform for machine learning.

In [3]:

```
import numpy as np
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense
from keras.activations import linear, relu, sigmoid
%matplotlib widget
import matplotlib.pyplot as plt
plt.style.use('./deeplearning.mplstyle')

import logging
logging.getLogger("tensorflow").setLevel(logging.ERROR)
tf.autograph.set_verbosity(0)

from public_tests import *

from autils import *
from lab_utils_softmax import plt_softmax
np.set_printoptions(precision=2)
```

2 - ReLU Activation

This week, a new activation was introduced, the Rectified Linear Unit (ReLU).

$$a = \max(0, z) \quad \# \text{ReLU function}$$

In [4]:

```
plt_act_trio()
```

The example from the lecture on the right shows an application of the ReLU. In this example, the derived "awareness" feature is not binary but has a continuous range of values. The sigmoid is best for on/off or binary situations. The ReLU provides a continuous linear relationship. Additionally it has an 'off' range where the output is zero.

The "off" feature makes the ReLU a Non-Linear activation. Why is this needed? This enables multiple units to contribute to the resulting function without interfering. This is examined more in the supporting optional lab.

3 - Softmax Function

A multiclass neural network generates N outputs. One output is selected as the predicted answer. In the output layer, a vector z is generated by a linear function which is fed into a softmax function. The softmax function converts z into a probability distribution as described below. After applying softmax, each output will be between 0 and 1 and the outputs will sum to 1. They can be interpreted as probabilities. The larger inputs to the softmax will correspond to larger output probabilities.

The softmax function can be written:

$$a_j = \frac{e^{z_j}}{\sum_{k=0}^{N-1} e^{z_k}}$$

Where $z = w \cdot x + b$

and N is the number of feature/categories in the output layer.

Exercise 1

Let's create a NumPy implementation:

In [5]:

```
# UNQ_C1
# GRADED CELL: my_softmax

def my_softmax(z):
    """ Softmax converts a vector of values to a probability distribution.
    Args:
        z (ndarray (N,)) : input data, N features
    Returns:
        a (ndarray (N,)) : softmax of z
    """
    ### START CODE HERE ###
    ez = np.exp(z)
    a = ez/np.sum(ez)
    ### END CODE HERE ###
    return a
```

In [6]:

```
z = np.array([1., 2., 3., 4.])
a = my_softmax(z)
atf = tf.nn.softmax(z)
```

```
print(f"my_softmax(z): {a}")
print(f"tensorflow softmax(z): {atf}")
```

```
# BEGIN UNIT TEST
test_my_softmax(my_softmax)
# END UNIT TEST
```

```
my_softmax(z): [0.03 0.09 0.24 0.64]
tensorflow softmax(z): [0.03 0.09 0.24 0.64]
All tests passed.
```

► Click for hints

Below, vary the values of the `z` inputs. Note in particular how the exponential in the numerator magnifies small differences in the values. Note as well that the output values sum to one.

In [7]:

```
plt.close("all")
plt_softmax(my_softmax)
```

4 - Neural Networks

In last weeks assignment, you implemented a neural network to do binary classification. This week you will extend that to multiclass classification. This will utilize the softmax activation.

4.1 Problem Statement

In this exercise, you will use a neural network to recognize ten handwritten digits, 0-9. This is a multiclass classification task where one of n choices is selected. Automated handwritten digit recognition is widely used today - from recognizing zip codes (postal codes) on mail envelopes to recognizing amounts written on bank checks.

4.2 Dataset

You will start by loading the dataset for this task.

- The `load_data()` function shown below loads the data into variables `X` and `y`
- The data set contains 5000 training examples of handwritten digits¹.
 - Each training example is a 20-pixel x 20-pixel grayscale image of the digit.
 - Each pixel is represented by a floating-point number indicating the grayscale intensity at that location.
 - The 20 by 20 grid of pixels is “unrolled” into a 400-dimensional vector.
 - Each training examples becomes a single row in our data matrix `X`.
 - This gives us a 5000 x 400 matrix `X` where every row is a training example of a handwritten digit image.

$$X = \begin{pmatrix} - & - & (x^{(1)}) & - & - \\ - & - & (x^{(2)}) & - & - \\ & & \vdots & & \\ - & - & (x^{(m)}) & - & - \end{pmatrix}$$

- The second part of the training set is a 5000 x 1 dimensional vector `y` that contains labels for the training set
 - `y = 0` if the image is of the digit `0`, `y = 4` if the image is of the digit `4` and so on.

In [8]:

```
# load dataset
X, y = load_data()
```

4.2.1 View the variables

Let's get more familiar with your dataset.

- A good place to start is to print out each variable and see what it contains.

The code below prints the first element in the variables `X` and `y`.

In [9]:

```
print ('The first element of X is: ', X[0])
```

The first element of X is: [0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00

0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	8.56e-06	1.94e-06	-7.37e-04
-8.13e-03	-1.86e-02	-1.87e-02	-1.88e-02	-1.91e-02	-1.64e-02	-3.78e-03
3.30e-04	1.28e-05	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	1.16e-04	1.20e-04	-1.40e-02	-2.85e-02	8.04e-02
2.67e-01	2.74e-01	2.79e-01	2.74e-01	2.25e-01	2.78e-02	-7.06e-03
2.35e-04	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
1.28e-17	-3.26e-04	-1.39e-02	8.16e-02	3.83e-01	8.58e-01	1.00e+00
9.70e-01	9.31e-01	1.00e+00	9.64e-01	4.49e-01	-5.60e-03	-3.78e-03
0.00e+00	0.00e+00	0.00e+00	0.00e+00	5.11e-06	4.36e-04	-3.96e-03
-2.69e-02	1.01e-01	6.42e-01	1.03e+00	8.51e-01	5.43e-01	3.43e-01
2.69e-01	6.68e-01	1.01e+00	9.04e-01	1.04e-01	-1.66e-02	0.00e+00
0.00e+00	0.00e+00	0.00e+00	2.60e-05	-3.11e-03	7.52e-03	1.78e-01
7.93e-01	9.66e-01	4.63e-01	6.92e-02	-3.64e-03	-4.12e-02	-5.02e-02
1.56e-01	9.02e-01	1.05e+00	1.51e-01	-2.16e-02	0.00e+00	0.00e+00
0.00e+00	5.87e-05	-6.41e-04	-3.23e-02	2.78e-01	9.37e-01	1.04e+00
5.98e-01	-3.59e-03	-2.17e-02	-4.81e-03	6.17e-05	-1.24e-02	1.55e-01
9.15e-01	9.20e-01	1.09e-01	-1.71e-02	0.00e+00	0.00e+00	1.56e-04
-4.28e-04	-2.51e-02	1.31e-01	7.82e-01	1.03e+00	7.57e-01	2.85e-01
4.87e-03	-3.19e-03	0.00e+00	8.36e-04	-3.71e-02	4.53e-01	1.03e+00
5.39e-01	-2.44e-03	-4.80e-03	0.00e+00	0.00e+00	-7.04e-04	-1.27e-02
1.62e-01	7.80e-01	1.04e+00	8.04e-01	1.61e-01	-1.38e-02	2.15e-03
-2.13e-04	2.04e-04	-6.86e-03	4.32e-04	7.21e-01	8.48e-01	1.51e-01
-2.28e-02	1.99e-04	0.00e+00	0.00e+00	-9.40e-03	3.75e-02	6.94e-01
1.03e+00	1.02e+00	8.80e-01	3.92e-01	-1.74e-02	-1.20e-04	5.55e-05
-2.24e-03	-2.76e-02	3.69e-01	9.36e-01	4.59e-01	-4.25e-02	1.17e-03
1.89e-05	0.00e+00	0.00e+00	-1.94e-02	1.30e-01	9.80e-01	9.42e-01
7.75e-01	8.74e-01	2.13e-01	-1.72e-02	0.00e+00	1.10e-03	-2.62e-02
1.23e-01	8.31e-01	7.27e-01	5.24e-02	-6.19e-03	0.00e+00	0.00e+00
0.00e+00	0.00e+00	-9.37e-03	3.68e-02	6.99e-01	1.00e+00	6.06e-01
3.27e-01	-3.22e-02	-4.83e-02	-4.34e-02	-5.75e-02	9.56e-02	7.27e-01
6.95e-01	1.47e-01	-1.20e-02	-3.03e-04	0.00e+00	0.00e+00	0.00e+00
0.00e+00	-6.77e-04	-6.51e-03	1.17e-01	4.22e-01	9.93e-01	8.82e-01
7.46e-01	7.24e-01	7.23e-01	7.20e-01	8.45e-01	8.32e-01	6.89e-02
-2.78e-02	3.59e-04	7.15e-05	0.00e+00	0.00e+00	0.00e+00	0.00e+00
1.53e-04	3.17e-04	-2.29e-02	-4.14e-03	3.87e-01	5.05e-01	7.75e-01
9.90e-01	1.01e+00	1.01e+00	7.38e-01	2.15e-01	-2.70e-02	1.33e-03
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	2.36e-04	-2.26e-03	-2.52e-02	-3.74e-02	6.62e-02	2.91e-01
3.23e-01	3.06e-01	8.76e-02	-2.51e-02	2.37e-04	0.00e+00	0.00e+00
0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00	0.00e+00
0.00e+00	0.00e+00	6.21e-18	6.73e-04	-1.13e-02	-3.55e-02	-3.88e-02

```
-3.71e-02 -1.34e-02 9.91e-04 4.89e-05 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00 0.00e+00
0.00e+00]
```

In [10]:

```
print ('The first element of y is: ', y[0,0])
print ('The last element of y is: ', y[-1,0])
```

```
The first element of y is: 0
The last element of y is: 9
```

4.2.2 Check the dimensions of your variables

Another way to get familiar with your data is to view its dimensions. Please print the shape of `X` and `y` and see how many training examples you have in your dataset.

In [11]:

```
print ('The shape of X is: ' + str(X.shape))
print ('The shape of y is: ' + str(y.shape))
```

```
The shape of X is: (5000, 400)
The shape of y is: (5000, 1)
```

4.2.3 Visualizing the Data

You will begin by visualizing a subset of the training set.

- In the cell below, the code randomly selects 64 rows from `X`, maps each row back to a 20 pixel by 20 pixel grayscale image and displays the images together.
- The label for each image is displayed above the image

In [12]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8,8, figsize=(5,5))
fig.tight_layout(pad=0.13,rect=[0, 0.03, 1, 0.91]) #[left, bottom, right, top]

#fig.tight_layout(pad=0.5)
widgvis(fig)
for i,ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)

    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_resaped = X[random_index].reshape((20,20)).T

    # Display the image
    ax.imshow(X_random_resaped, cmap='gray')

    # Display the label above the image
    ax.set_title(y[random_index,0])
    ax.set_axis_off()
    fig.suptitle("Label, image", fontsize=14)
```

4.3 Model representation

The neural network you will use in this assignment is shown in the figure below.

- This has two dense layers with ReLU activations followed by an output layer with a linear activation.
 - Recall that our inputs are pixel values of digit images.
 - Since the images are of size 28×28 , this gives us 784 inputs
- The parameters have dimensions that are sized for a neural network with 25 units in layer 1, 15 units in layer 2 and 10 output units in layer 3, one for each digit.
 - Recall that the dimensions of these parameters is determined as follows:
 - If network has s_{in} units in a layer and s_{out} units in the next layer, then
 - W will be of dimension $s_{in} \times s_{out}$.
 - b will be a vector with s_{out} elements
 - Therefore, the shapes of W , and b , are
 - layer1: The shape of W_1 is (784, 25) and the shape of b_1 is (25,)
 - layer2: The shape of W_2 is (25, 15) and the shape of b_2 is: (15,)
 - layer3: The shape of W_3 is (15, 10) and the shape of b_3 is: (10,)

Note: The bias vector b could be represented as a 1-D (n,) or 2-D (n,1) array. Tensorflow utilizes a 1-D representation and this lab will maintain that convention:

4.4 Tensorflow Model Implementation

Tensorflow models are built layer by layer. A layer's input dimensions (s_{in} above) are calculated for you. You specify a layer's *output dimensions* and this determines the next layer's input dimension. The input dimension of the first layer is derived from the size of the input data specified in the `model.fit` statement below.

Note: It is also possible to add an input layer that specifies the input dimension of the first layer. For example:

```
tf.keras.Input(shape=(784,)), #specify input shape
```

We will include that here to illuminate some model sizing.

4.5 Softmax placement

As described in the lecture and the optional softmax lab, numerical stability is improved if the softmax is grouped with the loss function rather than the output layer during training. This has implications when *building* the model and *using* the model.

Building:

- The final Dense layer should use a 'linear' activation. This is effectively no activation.

- The `model.compile` statement will indicate this by including `from_logits=True`.
`loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)`
- This does not impact the form of the target. In the case of `SparseCategoricalCrossentropy`, the target is the expected digit, 0-9.

Using the model:

- The outputs are not probabilities. If output probabilities are desired, apply a softmax function.

Exercise 2

Below, using Keras [Sequential model](#) and [Dense Layer](#) with a ReLU activation to construct the three layer network described above.

In [13]:

```
# UNQ_C2
# GRADED CELL: Sequential model
tf.random.set_seed(1234) # for consistent results
model = Sequential(
    [
        ### START CODE HERE ###
        tf.keras.layers.InputLayer((400,)),
        tf.keras.layers.Dense(25, activation="relu", name = "L1"),
        tf.keras.layers.Dense(15, activation="relu", name = "L2"),
        tf.keras.layers.Dense(10, activation="linear", name = "L3"),
        ### END CODE HERE ###
    ], name = "my_model"
)
```

In [14]:

```
model.summary()
```

Model: "my_model"

Layer (type)	Output Shape	Param #
=====		
L1 (Dense)	(None, 25)	10025
L2 (Dense)	(None, 15)	390
L3 (Dense)	(None, 10)	160
=====		
Total params: 10575 (41.31 KB)		
Trainable params: 10575 (41.31 KB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

► **Expected Output (Click to expand)**

► **Click for hints**

In [15]:

```
# BEGIN UNIT TEST
test_model(model, 10, 400)
# END UNIT TEST
```

All tests passed!

The parameter counts shown in the summary correspond to the number of elements in the weight and bias arrays as shown below.

Let's further examine the weights to verify that tensorflow produced the same dimensions as we calculated above.

In [16]:

```
[layer1, layer2, layer3] = model.layers
```

In [17]:

```
#### Examine Weights shapes
W1,b1 = layer1.get_weights()
W2,b2 = layer2.get_weights()
W3,b3 = layer3.get_weights()
print(f"W1 shape = {W1.shape}, b1 shape = {b1.shape}")
print(f"W2 shape = {W2.shape}, b2 shape = {b2.shape}")
print(f"W3 shape = {W3.shape}, b3 shape = {b3.shape}")
```

```
W1 shape = (400, 25), b1 shape = (25,)
W2 shape = (25, 15), b2 shape = (15,)
W3 shape = (15, 10), b3 shape = (10,)
```

Expected Output

```
W1 shape = (400, 25), b1 shape = (25,)
W2 shape = (25, 15), b2 shape = (15,)
W3 shape = (15, 10), b3 shape = (10,)
```

The following code:

- defines a loss function, `SparseCategoricalCrossentropy` and indicates the softmax should be included with the loss calculation by adding `from_logits=True`
- defines an optimizer. A popular choice is Adaptive Moment (Adam) which was described in lecture.

In [18]:

```
model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
)

history = model.fit(
    X,y,
    epochs=40
)
```

```
Epoch 1/40
157/157 [=====] - 1s 1ms/step - loss: 1.5852
Epoch 2/40
157/157 [=====] - 0s 1ms/step - loss: 0.5685
Epoch 3/40
157/157 [=====] - 0s 1ms/step - loss: 0.3881
Epoch 4/40
157/157 [=====] - 0s 1ms/step - loss: 0.3155
Epoch 5/40
157/157 [=====] - 0s 1ms/step - loss: 0.2773
Epoch 6/40
157/157 [=====] - 0s 1ms/step - loss: 0.2494
Epoch 7/40
157/157 [=====] - 0s 999us/step - loss: 0.2289
Epoch 8/40
157/157 [=====] - 0s 1ms/step - loss: 0.2101
Epoch 9/40
157/157 [=====] - 0s 992us/step - loss: 0.1978
Epoch 10/40
157/157 [=====] - 0s 988us/step - loss: 0.1833
Epoch 11/40
157/157 [=====] - 0s 977us/step - loss: 0.1699
Epoch 12/40
```



```

Epoch 12/40
157/157 [=====] - 0s 972us/step - loss: 0.1608
Epoch 13/40
157/157 [=====] - 0s 990us/step - loss: 0.1503
Epoch 14/40
157/157 [=====] - 0s 996us/step - loss: 0.1437
Epoch 15/40
157/157 [=====] - 0s 959us/step - loss: 0.1338
Epoch 16/40
157/157 [=====] - 0s 977us/step - loss: 0.1288
Epoch 17/40
157/157 [=====] - 0s 1ms/step - loss: 0.1207
Epoch 18/40
157/157 [=====] - 0s 977us/step - loss: 0.1138
Epoch 19/40
157/157 [=====] - 0s 980us/step - loss: 0.1073
Epoch 20/40
157/157 [=====] - 0s 969us/step - loss: 0.1035
Epoch 21/40
157/157 [=====] - 0s 979us/step - loss: 0.0967
Epoch 22/40
157/157 [=====] - 0s 985us/step - loss: 0.0925
Epoch 23/40
157/157 [=====] - 0s 968us/step - loss: 0.0868
Epoch 24/40
157/157 [=====] - 0s 982us/step - loss: 0.0796
Epoch 25/40
157/157 [=====] - 0s 978us/step - loss: 0.0778
Epoch 26/40
157/157 [=====] - 0s 984us/step - loss: 0.0737
Epoch 27/40
157/157 [=====] - 0s 987us/step - loss: 0.0710
Epoch 28/40
157/157 [=====] - 0s 970us/step - loss: 0.0651
Epoch 29/40
157/157 [=====] - 0s 976us/step - loss: 0.0606
Epoch 30/40
157/157 [=====] - 0s 1ms/step - loss: 0.0603
Epoch 31/40
157/157 [=====] - 0s 966us/step - loss: 0.0553
Epoch 32/40
157/157 [=====] - 0s 975us/step - loss: 0.0511
Epoch 33/40
157/157 [=====] - 0s 956us/step - loss: 0.0471
Epoch 34/40
157/157 [=====] - 0s 969us/step - loss: 0.0445
Epoch 35/40
157/157 [=====] - 0s 963us/step - loss: 0.0426
Epoch 36/40
157/157 [=====] - 0s 962us/step - loss: 0.0414
Epoch 37/40
157/157 [=====] - 0s 973us/step - loss: 0.0372
Epoch 38/40
157/157 [=====] - 0s 964us/step - loss: 0.0371
Epoch 39/40
157/157 [=====] - 0s 968us/step - loss: 0.0337
Epoch 40/40
157/157 [=====] - 0s 985us/step - loss: 0.0304

```

Epochs and batches

In the `compile` statement above, the number of `epochs` was set to 100. This specifies that the entire data set should be applied during training 100 times. During training, you see output describing the progress of training that looks like this:

```

Epoch 1/100
157/157 [=====] - 0s 1ms/step - loss: 2.2770

```

The first line, `Epoch 1/100`, describes which epoch the model is currently running. For efficiency, the training data set is broken into 'batches'. The default size of a batch in Tensorflow is 32. There are 5000 examples in our

data set or roughly 157 batches. The notation on the 2nd line 157/157 [====] is describing which batch has been executed.

Loss (cost)

In course 1, we learned to track the progress of gradient descent by monitoring the cost. Ideally, the cost will decrease as the number of iterations of the algorithm increases. Tensorflow refers to the cost as `loss`. Above, you saw the loss displayed each epoch as `model.fit` was executing. The `.fit` method returns a variety of metrics including the loss. This is captured in the `history` variable above. This can be used to examine the loss in a plot as shown below.

In [19]:

```
plot_loss_tf(history)
```

Prediction

To make a prediction, use Keras `predict`. Below, `X[1015]` contains an image of a two.

In [20]:

```
image_of_two = X[1015]
display_digit(image_of_two)

prediction = model.predict(image_of_two.reshape(1,400)) # prediction

print(f" predicting a Two: \n{prediction}")
print(f" Largest Prediction index: {np.argmax(prediction)}")
```

```
1/1 [=====] - 0s 69ms/step
predicting a Two:
[[-6.46  1.66  3.11  0.94 -12.55 -8.78 -5.07 -1.3  -4.94 -7.84]]
Largest Prediction index: 2
```

The largest output is `prediction[2]`, indicating the predicted digit is a '2'. If the problem only requires a selection, that is sufficient. Use NumPy [argmax](#) to select it. If the problem requires a probability, a softmax is required:

In [21]:

```
prediction_p = tf.nn.softmax(prediction)

print(f" predicting a Two. Probability vector: \n{prediction_p}")
print(f"Total of predictions: {np.sum(prediction_p):0.3f}")

predicting a Two. Probability vector:
[[5.16e-05 1.72e-01 7.35e-01 8.38e-02 1.16e-07 5.05e-06 2.06e-04 8.96e-03
 2.34e-04 1.29e-05]]
Total of predictions: 1.000
```

To return an integer representing the predicted target, you want the index of the largest probability. This is accomplished with the Numpy [argmax](#) function.

In [22]:

```
yhat = np.argmax(prediction_p)

print(f"np.argmax(prediction_p): {yhat}")

np.argmax(prediction_p): 2
```

Let's compare the predictions vs the labels for a random sample of 64 digits. This takes a moment to run.

In [23]:

```

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8,8, figsize=(5,5))
fig.tight_layout(pad=0.13,rect=[0, 0.03, 1, 0.91]) #[left, bottom, right, top]
widgvis(fig)
for i,ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)

    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_resaped = X[random_index].reshape((20,20)).T

    # Display the image
    ax.imshow(X_random_resaped, cmap='gray')

    # Predict using the Neural Network
    prediction = model.predict(X[random_index].reshape(1,400))
    prediction_p = tf.nn.softmax(prediction)
    yhat = np.argmax(prediction_p)

    # Display the label above the image
    ax.set_title(f"{y[random_index,0]}, {yhat}", fontsize=10)
    ax.set_axis_off()
fig.suptitle("Label, yhat", fontsize=14)
plt.show()

```

```

1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 14ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step

```

```

1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 14ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 14ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step

```

Let's look at some of the errors.

Note: increasing the number of training epochs can eliminate the errors on this data set.

In [24]:

```
print( f"{display_errors(model,X,y)} errors out of {len(X)} images")
```

```

157/157 [=====] - 0s 801us/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
29 errors out of 5000 images

```

Congratulations!

You have successfully built and utilized a neural network to do multiclass classification.

► [Please click here if you want to experiment with any of the non-graded code.](#)