

Practice Lab: Neural Networks for Handwritten Digit Recognition, Binary

In this exercise, you will use a neural network to recognize the hand-written digits zero and one.

Outline

- [1 - Packages](#)
- [2 - Neural Networks](#)
 - [2.1 Problem Statement](#)
 - [2.2 Dataset](#)
 - [2.3 Model representation](#)
 - [2.4 Tensorflow Model Implementation](#)
 - [Exercise 1](#)
 - [2.5 NumPy Model Implementation \(Forward Prop in NumPy\)](#)
 - [Exercise 2](#)
 - [2.6 Vectorized NumPy Model Implementation \(Optional\)](#)
 - [Exercise 3](#)
 - [2.7 Congratulations!](#)
 - [2.8 NumPy Broadcasting Tutorial \(Optional\)](#)

NOTE: To prevent errors from the autograder, you are not allowed to edit or delete non-graded cells in this notebook . Please also refrain from adding any new cells. **Once you have passed this assignment** and want to experiment with any of the non-graded code, you may follow the instructions at the bottom of this notebook.

1 - Packages

First, let's run the cell below to import all the packages that you will need during this assignment.

- [numpy](#) is the fundamental package for scientific computing with Python.
- [matplotlib](#) is a popular library to plot graphs in Python.
- [tensorflow](#) a popular platform for machine learning.

In [7]:

```
import numpy as np
import tensorflow as tf
import keras
from keras.models import Sequential
from keras.layers import Dense
import matplotlib.pyplot as plt
from autils import *
%matplotlib inline

import logging
logging.getLogger("tensorflow").setLevel(logging.ERROR)
tf.autograph.set_verbosity(0)
```

Tensorflow and Keras

Tensorflow is a machine learning package developed by Google. In 2019, Google integrated Keras into Tensorflow and released Tensorflow 2.0. Keras is a framework developed independently by François Chollet that creates a simple, layer-centric interface to Tensorflow. This course will be using the Keras interface.

2 - Neural Networks

In Course 1, you implemented logistic regression. This was extended to handle non-linear boundaries using polynomial regression. For even more complex scenarios such as image recognition, neural networks are preferred.

2.1 Problem Statement

In this exercise, you will use a neural network to recognize two handwritten digits, zero and one. This is a binary classification task. Automated handwritten digit recognition is widely used today - from recognizing zip codes (postal codes) on mail envelopes to recognizing amounts written on bank checks. You will extend this network to recognize all 10 digits (0-9) in a future assignment.

This exercise will show you how the methods you have learned can be used for this classification task.

2.2 Dataset

You will start by loading the dataset for this task.

- The `load_data()` function shown below loads the data into variables `X` and `y`
- The data set contains 1000 training examples of handwritten digits ¹, here limited to zero and one.
 - Each training example is a 20-pixel x 20-pixel grayscale image of the digit.
 - Each pixel is represented by a floating-point number indicating the grayscale intensity at that location.
 - The 20 by 20 grid of pixels is “unrolled” into a 400-dimensional vector.
 - Each training example becomes a single row in our data matrix `X`.
 - This gives us a 1000 x 400 matrix `X` where every row is a training example of a handwritten digit image.

$$X = \begin{pmatrix} \text{---} (x^{(1)}) \text{---} \\ \text{---} (x^{(2)}) \text{---} \\ \vdots \\ \text{---} (x^{(m)}) \text{---} \end{pmatrix}$$

- The second part of the training set is a 1000 x 1 dimensional vector `y` that contains labels for the training set
 - `y = 0` if the image is of the digit `0`, `y = 1` if the image is of the digit `1`.

¹ This is a subset of the MNIST handwritten digit dataset (<http://yann.lecun.com/exdb/mnist/>)

In [8]:

```
# load dataset
X, y = load_data()
```

2.2.1 View the variables

Let's get more familiar with your dataset.

- A good place to start is to print out each variable and see what it contains.

The code below prints elements of the variables `X` and `y`.

In [9]:

```
print ('The first element of X is: ', X[0])
```

```
The first element of X is: [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
```

[illegible]

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8,8, figsize=(8,8))
fig.tight_layout(pad=0.1)
```

```


for i, ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)

    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_resized = X[random_index].reshape((20,20)).T

    # Display the image
    ax.imshow(X_random_resized, cmap='gray')

    # Display the label above the image
    ax.set_title(y[random_index,0])
    ax.set_axis_off()

```



2.3 Model representation

The neural network you will use in this assignment is shown in the figure below.

- This has three dense layers with sigmoid activations.
 - Recall that our inputs are pixel values of digit images.

- Since the images are of size 20×20 , this gives us 400 inputs

- The parameters have dimensions that are sized for a neural network with 25 units in layer 1, 15 units in layer 2 and 1 output unit in layer 3.
 - Recall that the dimensions of these parameters are determined as follows:
 - If network has s_{in} units in a layer and s_{out} units in the next layer, then
 - W will be of dimension $s_{in} \times s_{out}$
 - b will be a vector with s_{out} elements
 - Therefore, the shapes of W , and b , are
 - layer1: The shape of W_1 is (400, 25) and the shape of b_1 is (25,)
 - layer2: The shape of W_2 is (25, 15) and the shape of b_2 is: (15,)
 - layer3: The shape of W_3 is (15, 1) and the shape of b_3 is: (1,)

Note: The bias vector b could be represented as a 1-D (n,) or 2-D (1,n) array. Tensorflow utilizes a 1-D representation and this lab will maintain that convention.

2.4 Tensorflow Model Implementation

Tensorflow models are built layer by layer. A layer's input dimensions (s_{in} above) are calculated for you. You specify a layer's *output dimensions* and this determines the next layer's input dimension. The input dimension of the first layer is derived from the size of the input data specified in the `model.fit` statement below.

Note: It is also possible to add an input layer that specifies the input dimension of the first layer. For example:

```
tf.keras.Input(shape=(400,)), #specify input shape
```

We will include that here to illuminate some model sizing.

Exercise 1

Below, using Keras [Sequential model](#) and [Dense Layer](#) with a sigmoid activation to construct the network described above.

In [15]:

```
# UNQ_C1
# GRADED CELL: Sequential model

model = Sequential(
    [
        tf.keras.Input(shape=(400,)), #specify input size
        ### START CODE HERE ###
        tf.keras.layers.Dense(25, activation="sigmoid"),
        tf.keras.layers.Dense(15, activation="sigmoid"),
        tf.keras.layers.Dense(1, activation="sigmoid")
        ### END CODE HERE ###
    ], name = "my_model"
)
```

In [16]:

```
model.summary()
```

Model: "my_model"

Layer (type)	Output Shape	Param #
Input Layer	(400,)	0
Dense (25)	(25,)	10000
Dense (15)	(15,)	375
Dense (1)	(1,)	15
Total		10390

dense (Dense)	(None, 25)	10025
dense_1 (Dense)	(None, 15)	390
dense_2 (Dense)	(None, 1)	16

```
=====
Total params: 10431 (40.75 KB)
Trainable params: 10431 (40.75 KB)
Non-trainable params: 0 (0.00 Byte)
```

► Expected Output (Click to Expand)

► Click for hints

In [17]:

```
# UNIT TESTS
from public_tests import *

test_c1(model)
```

All tests passed!

The parameter counts shown in the summary correspond to the number of elements in the weight and bias arrays as shown below.

In [18]:

```
L1_num_params = 400 * 25 + 25 # W1 parameters + b1 parameters
L2_num_params = 25 * 15 + 15  # W2 parameters + b2 parameters
L3_num_params = 15 * 1 + 1    # W3 parameters + b3 parameters
print("L1 params = ", L1_num_params, ", L2 params = ", L2_num_params, ", L3 params = ",
      L3_num_params )
```

L1 params = 10025 , L2 params = 390 , L3 params = 16

We can examine details of the model by first extracting the layers with `model.layers` and then extracting the weights with `layerx.get_weights()` as shown below.

In [19]:

```
[layer1, layer2, layer3] = model.layers
```

In [20]:

```
#### Examine Weights shapes
W1,b1 = layer1.get_weights()
W2,b2 = layer2.get_weights()
W3,b3 = layer3.get_weights()
print(f"W1 shape = {W1.shape}, b1 shape = {b1.shape}")
print(f"W2 shape = {W2.shape}, b2 shape = {b2.shape}")
print(f"W3 shape = {W3.shape}, b3 shape = {b3.shape}")
```

```
W1 shape = (400, 25), b1 shape = (25,)
W2 shape = (25, 15), b2 shape = (15,)
W3 shape = (15, 1), b3 shape = (1,)
```

Expected Output

```
W1 shape = (400, 25), b1 shape = (25,)
W2 shape = (25, 15), b2 shape = (15,)
W3 shape = (15, 1), b3 shape = (1,)
```

`xx.get_weights` returns a NumPy array. One can also access the weights directly in their tensor form. Note

the shape of the tensors in the final layer.

In [21]:

```
print(model.layers[2].weights)
```

```
[<tf.Variable 'dense_2/kernel:0' shape=(15, 1) dtype=float32, numpy=
array([[ -0.1633364 ],
       [  0.5309529 ],
       [  0.09172976],
       [-0.30924553],
       [-0.18748146],
       [-0.5079004 ],
       [  0.08098614],
       [-0.201843  ],
       [-0.27523562],
       [-0.27439743],
       [-0.21130604],
       [-0.4071542 ],
       [  0.12160844],
       [  0.32570118],
       [-0.29153123]], dtype=float32)>, <tf.Variable 'dense_2/bias:0' shape=(1,) dtype=fl
oat32, numpy=array([0.], dtype=float32)>]
```

The following code will define a loss function and run gradient descent to fit the weights of the model to the training data. This will be explained in more detail in the following week.

In [22]:

```
model.compile(
    loss=tf.keras.losses.BinaryCrossentropy(),
    optimizer=tf.keras.optimizers.Adam(0.001),
)

model.fit(
    X,y,
    epochs=20
)
```

```
Epoch 1/20
32/32 [=====] - 1s 1ms/step - loss: 0.6691
Epoch 2/20
32/32 [=====] - 0s 1ms/step - loss: 0.5326
Epoch 3/20
32/32 [=====] - 0s 1ms/step - loss: 0.4017
Epoch 4/20
32/32 [=====] - 0s 1ms/step - loss: 0.2932
Epoch 5/20
32/32 [=====] - 0s 1ms/step - loss: 0.2182
Epoch 6/20
32/32 [=====] - 0s 1ms/step - loss: 0.1691
Epoch 7/20
32/32 [=====] - 0s 1ms/step - loss: 0.1362
Epoch 8/20
32/32 [=====] - 0s 1ms/step - loss: 0.1129
Epoch 9/20
32/32 [=====] - 0s 1ms/step - loss: 0.0958
Epoch 10/20
32/32 [=====] - 0s 1ms/step - loss: 0.0824
Epoch 11/20
32/32 [=====] - 0s 1ms/step - loss: 0.0722
Epoch 12/20
32/32 [=====] - 0s 1ms/step - loss: 0.0639
Epoch 13/20
32/32 [=====] - 0s 1ms/step - loss: 0.0572
Epoch 14/20
32/32 [=====] - 0s 1ms/step - loss: 0.0516
Epoch 15/20
32/32 [=====] - 0s 1ms/step - loss: 0.0470
Epoch 16/20
32/32 [=====] - 0s 1ms/step - loss: 0.0430
```



```
Epoch 17/20
32/32 [=====] - 0s 1ms/step - loss: 0.0396
Epoch 18/20
32/32 [=====] - 0s 1ms/step - loss: 0.0367
Epoch 19/20
32/32 [=====] - 0s 1ms/step - loss: 0.0342
Epoch 20/20
32/32 [=====] - 0s 1ms/step - loss: 0.0320
```

Out[22]:

```
<keras.src.callbacks.History at 0x1fb691e3450>
```

To run the model on an example to make a prediction, use `Keras predict`. The input to `predict` is an array so the single example is reshaped to be two dimensional.

In [23]:

```
prediction = model.predict(X[0].reshape(1,400)) # a zero
print(f" predicting a zero: {prediction}")
prediction = model.predict(X[500].reshape(1,400)) # a one
print(f" predicting a one: {prediction}")

1/1 [=====] - 0s 73ms/step
predicting a zero: [[0.01334678]]
1/1 [=====] - 0s 19ms/step
predicting a one: [[0.9608029]]
```

The output of the model is interpreted as a probability. In the first example above, the input is a zero. The model predicts the probability that the input is a one is nearly zero. In the second example, the input is a one. The model predicts the probability that the input is a one is nearly one. As in the case of logistic regression, the probability is compared to a threshold to make a final prediction.

In [24]:

```
if prediction >= 0.5:
    yhat = 1
else:
    yhat = 0
print(f"prediction after threshold: {yhat}")
```

```
prediction after threshold: 1
```

Let's compare the predictions vs the labels for a random sample of 64 digits. This takes a moment to run.

In [25]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8,8, figsize=(8,8))
fig.tight_layout(pad=0.1,rect=[0, 0.03, 1, 0.92]) #[left, bottom, right, top]

for i,ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)

    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_reshaped = X[random_index].reshape((20,20)).T

    # Display the image
    ax.imshow(X_random_reshaped, cmap='gray')

    # Predict using the Neural Network
    prediction = model.predict(X[random_index].reshape(1,400))
```

```

if prediction >= 0.5:
    yhat = 1
else:
    yhat = 0

# Display the label above the image
ax.set_title(f"{y[random_index,0]},{yhat}")
ax.set_axis_off()
fig.suptitle("Label, yhat", fontsize=16)
plt.show()

```

```

1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step


```

```

1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step

```

Label, yhat



2.5 NumPy Model Implementation (Forward Prop in NumPy)

As described in lecture, it is possible to build your own dense layer using NumPy. This can then be utilized to build a multi-layer neural network.

Exercise 2

Below, build a dense layer subroutine. The example in lecture utilized a for loop to visit each unit (j) in the layer and perform the dot product of the weights for that unit ($w[:,j]$) and sum the bias for the unit ($b[j]$) to form z . An activation function $g(z)$ is then applied to that result. This section will not utilize some of the matrix operations described in the optional lectures. These will be explored in a later section.

In [28]:

```
# UNQ_C2
```

```
# GRADED FUNCTION: my_dense
```

```
def my_dense(a_in, W, b, g):  
    """  
    Computes dense layer  
    Args:  
        a_in (ndarray (n, )) : Data, 1 example  
        W     (ndarray (n,j)) : Weight matrix, n features per unit, j units  
        b     (ndarray (j, )) : bias vector, j units  
        g     activation function (e.g. sigmoid, relu..)  
    Returns  
        a_out (ndarray (j,)) : j units  
    """  
    units = W.shape[1]  
    a_out = np.zeros(units)  
    ### START CODE HERE ###  
    for i in range (units):  
        w = W[:,i]  
        z = np.dot(w,a_in) + b[i]  
        a_out[i] = g(z)  
    ### END CODE HERE ###  
    return (a_out)
```

In [29]:

```
# Quick Check  
x_tst = 0.1*np.arange(1,3,1).reshape(2,) # (1 examples, 2 features)  
W_tst = 0.1*np.arange(1,7,1).reshape(2,3) # (2 input features, 3 output features)  
b_tst = 0.1*np.arange(1,4,1).reshape(3,) # (3 features)  
A_tst = my_dense(x_tst, W_tst, b_tst, sigmoid)  
print(A_tst)
```

```
[0.54735762 0.57932425 0.61063923]
```

Expected Output

```
[0.54735762 0.57932425 0.61063923]
```

► Click for hints

In [30]:

```
# UNIT TESTS
```

```
test_c2(my_dense)
```

```
All tests passed!
```

The following cell builds a three-layer neural network utilizing the `my_dense` subroutine above.

In [31]:

```
def my_sequential(x, W1, b1, W2, b2, W3, b3):  
    a1 = my_dense(x, W1, b1, sigmoid)  
    a2 = my_dense(a1, W2, b2, sigmoid)  
    a3 = my_dense(a2, W3, b3, sigmoid)  
    return (a3)
```

We can copy trained weights and biases from Tensorflow.

In [32]:

```
W1_tmp,b1_tmp = layer1.get_weights()  
W2_tmp,b2_tmp = layer2.get_weights()  
W3_tmp,b3_tmp = layer3.get_weights()
```

In [33]:

```
# make predictions
prediction = my_sequential(X[0], W1_tmp, b1_tmp, W2_tmp, b2_tmp, W3_tmp, b3_tmp )
if prediction >= 0.5:
    yhat = 1
else:
    yhat = 0
print( "yhat = ", yhat, " label= ", y[0,0])
prediction = my_sequential(X[500], W1_tmp, b1_tmp, W2_tmp, b2_tmp, W3_tmp, b3_tmp )
if prediction >= 0.5:
    yhat = 1
else:
    yhat = 0
print( "yhat = ", yhat, " label= ", y[500,0])
```

```
yhat = 0  label= 0
yhat = 1  label= 1
```

Run the following cell to see predictions from both the Numpy model and the Tensorflow model. This takes a moment to run.

In [34]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8,8, figsize=(8,8))
fig.tight_layout(pad=0.1,rect=[0, 0.03, 1, 0.92]) #[left, bottom, right, top]

for i,ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)

    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_resaped = X[random_index].reshape((20,20)).T

    # Display the image
    ax.imshow(X_random_resaped, cmap='gray')

    # Predict using the Neural Network implemented in Numpy
    my_prediction = my_sequential(X[random_index], W1_tmp, b1_tmp, W2_tmp, b2_tmp, W3_tm
p, b3_tmp )
    my_yhat = int(my_prediction >= 0.5)

    # Predict using the Neural Network implemented in Tensorflow
    tf_prediction = model.predict(X[random_index].reshape(1,400))
    tf_yhat = int(tf_prediction >= 0.5)

    # Display the label above the image
    ax.set_title(f"{y[random_index,0]},{tf_yhat},{my_yhat}")
    ax.set_axis_off()
fig.suptitle("Label, yhat Tensorflow, yhat Numpy", fontsize=16)
plt.show()
```

```
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
```

C:\Users\fatih\AppData\Local\Temp\ipykernel_14404\1208669724.py:23: DeprecationWarning: C onversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecatcd NumPy 1.25.)

```
my_yhat = int(my_prediction >= 0.5)
```


C:\Users\fatih\AppData\Local\Temp\ipykernel_14404\1208669724.py:27: DeprecationWarning: C onversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecatcd NumPy 1.25.)


```
tf_yhat = int(tf_prediction >= 0.5)
```

ci_yhat = int(ci_prediction >= 0.5)

```
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
```

Label, yhat Tensorflow, yhat Numpy





2.6 Vectorized NumPy Model Implementation (Optional)

The optional lectures described vector and matrix operations that can be used to speed the calculations. Below describes a layer operation that computes the output for all units in a layer on a given input example:

We can demonstrate this using the examples `x` and the `W1`, `b1` parameters above. We use `np.matmul` to perform the matrix multiply. Note, the dimensions of `x` and `W` must be compatible as shown in the diagram above.

In [35]:

```
x = X[0].reshape(-1,1)           # column vector (400,1)
z1 = np.matmul(x.T,W1) + b1      # (1,400)(400,25) = (1,25)
a1 = sigmoid(z1)
print(a1.shape)
```

(1, 25)

You can take this a step further and compute all the units for all examples in one Matrix-Matrix operation.

The full operation is $Z = XW + b$

. This will utilize NumPy broadcasting to expand `b` to `m`

rows. If this is unfamiliar, a short tutorial is provided at the end of the notebook.

Exercise 3

Below, compose a new `my_dense_v` subroutine that performs the layer calculations for a matrix of examples.

This will utilize `np.matmul()` .

Note: This function is not graded because it is discussed in the optional lectures on vectorization. If you didn't go through them, feel free to click the hints below the expected code to see the code. You can also submit the notebook even with a blank answer here.

In [36]:

```
# UNQ_C3
# UNGRADED FUNCTION: my_dense_v

def my_dense_v(A_in, W, b, g):
    """
    Computes dense layer
    Args:
        A_in (ndarray (m,n)) : Data, m examples, n features each
        W      (ndarray (n,j)) : Weight matrix, n features per unit, j units
        b      (ndarray (1,j)) : bias vector, j units
        g      activation function (e.g. sigmoid, relu..)
    Returns
        A_out (tf.Tensor or ndarray (m,j)) : m examples, j units
    """
    ### START CODE HERE ###
    A_out = g(np.matmul(A_in,W) + b)
    ### END CODE HERE ###
    return(A_out)
```

In [37]:

```
X_tst = 0.1*np.arange(1,9,1).reshape(4,2) # (4 examples, 2 features)
W_tst = 0.1*np.arange(1,7,1).reshape(2,3) # (2 input features, 3 output features)
b_tst = 0.1*np.arange(1,4,1).reshape(1,3) # (1,3 features)
A_tst = my_dense_v(X_tst, W_tst, b_tst, sigmoid)
print(A_tst)
```

```
[[0.54735762 0.57932425 0.61063923]
 [0.57199613 0.61301418 0.65248946]
 [0.5962827  0.64565631 0.6921095 ]
 [0.62010643 0.67699586 0.72908792]]
```

Expected Output

```
[[0.54735762 0.57932425 0.61063923]
 [0.57199613 0.61301418 0.65248946]
 [0.5962827  0.64565631 0.6921095 ]
 [0.62010643 0.67699586 0.72908792]]
```

► Click for hints

In [38]:

```
# UNIT TESTS

test_c3(my_dense_v)
```

All tests passed!

The following cell builds a three-layer neural network utilizing the `my_dense_v` subroutine above.

In [39]:

```
def my_sequential_v(X, W1, b1, W2, b2, W3, b3):
    A1 = my_dense_v(X, W1, b1, sigmoid)
    A2 = my_dense_v(A1, W2, b2, sigmoid)
    A3 = my_dense_v(A2, W3, b3, sigmoid)
    return(A3)
```


We can again copy trained weights and biases from Tensorflow.

In [40]:

```
W1_tmp,b1_tmp = layer1.get_weights()
W2_tmp,b2_tmp = layer2.get_weights()
W3_tmp,b3_tmp = layer3.get_weights()
```

Let's make a prediction with the new model. This will make a prediction on *all of the examples at once*. Note the shape of the output.

In [41]:

```
Prediction = my_sequential_v(X, W1_tmp, b1_tmp, W2_tmp, b2_tmp, W3_tmp, b3_tmp )
Prediction.shape
```

Out[41]:

```
(1000, 1)
```

We'll apply a threshold of 0.5 as before, but to all predictions at once.

In [42]:

```
Yhat = (Prediction >= 0.5).astype(int)
print("predict a zero: ",Yhat[0], "predict a one: ", Yhat[500])
```

```
predict a zero:  [0] predict a one:  [1]
```

Run the following cell to see predictions. This will use the predictions we just calculated above. This takes a moment to run.

In [52]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# You do not need to modify anything in this cell

m, n = X.shape

fig, axes = plt.subplots(8, 8, figsize=(8, 8))
fig.tight_layout(pad=0.1, rect=[0, 0.03, 1, 0.92]) #[left, bottom, right, top]

for i, ax in enumerate(axes.flat):
    # Select random indices
    random_index = np.random.randint(m)


    # Select rows corresponding to the random indices and
    # reshape the image
    X_random_resaped = X[random_index].reshape((20, 20)).T

    # Display the image
    ax.imshow(X_random_resaped, cmap='gray')

    # Display the label above the image
    ax.set_title(f"{Y[random_index,0]}, {Yhat[random_index, 0]}")
    ax.set_axis_off()
fig.suptitle("Label, Yhat", fontsize=16)
plt.show()
```

Label, Yhat





You can see how one of the misclassified images looks.

In [53]:

```
fig = plt.figure(figsize=(1, 1))
errors = np.where(y != Yhat)
random_index = errors[0][0]
X_random_resaped = X[random_index].reshape((20, 20)).T
plt.imshow(X_random_resaped, cmap='gray')
plt.title(f"{y[random_index,0]}, {Yhat[random_index, 0]}")
plt.axis('off')
plt.show()
```

0, 1



2.7 Congratulations!

You have successfully built and utilized a neural network.

2.8 NumPy Broadcasting Tutorial (Optional)

In the last example, $Z = XW + b$

utilized NumPy broadcasting to expand the vector b

If you are not familiar with NumPy Broadcasting, this short tutorial is provided

. If you are not familiar with NumPy Broadcasting, this short tutorial is provided.

XW

is a matrix-matrix operation with dimensions $(m, j_1)(j_1, j_2)$

which results in a matrix with dimension (m, j_2)

. To that, we add a vector b

with dimension $(1, j_2)$

. b

must be expanded to be a (m, j_2)

matrix for this element-wise operation to make sense. This expansion is accomplished for you by NumPy broadcasting.

Broadcasting applies to element-wise operations.

Its basic operation is to 'stretch' a smaller dimension by replicating elements to match a larger dimension.

More [specifically](#): When operating on two arrays, NumPy compares their shapes element-wise. It starts with the trailing (i.e. rightmost) dimensions and works its way left. Two dimensions are compatible when

- they are equal, or
- one of them is 1

If these conditions are not met, a `ValueError: operands could not be broadcast together` exception is thrown, indicating that the arrays have incompatible shapes. The size of the resulting array is the size that is not 1 along each axis of the inputs.

Here are some examples:

missing

Calculating Broadcast Result shape

The graphic below describes expanding dimensions. Note the red text below:

missing

Broadcast notionally expands arguments to match for element wise operations

The graphic above shows NumPy expanding the arguments to match before the final operation. Note that this is a notional description. The actual mechanics of NumPy operation choose the most efficient implementation.

For each of the following examples, try to guess the size of the result before running the example.

In [54]:

```
a = np.array([1,2,3]).reshape(-1,1)  #(3,1)
b = 5
print(f"(a + b).shape: {(a + b).shape}, \na + b = \n{a + b}")

(a + b).shape: (3, 1),
a + b =
[[6]
 [7]
 [8]]
```

Note that this applies to all element-wise operations:

In [55]:

```
a = np.array([1,2,3]).reshape(-1,1)  #(3,1)
b = 5
print(f"(a * b).shape: {(a * b).shape}, \na * b = \n{a * b}")

(a * b).shape: (3, 1),
a * b =
[[ 5]
```

```
[10]  
[15]]
```

missing

Row-Column Element-Wise Operations

In [56]:

```
a = np.array([1,2,3,4]).reshape(-1,1)  
b = np.array([1,2,3]).reshape(1,-1)  
print(a)  
print(b)  
print(f"(a + b).shape: {(a + b).shape}, \na + b = \n{a + b}")
```

```
[[1]  
 [2]  
 [3]  
 [4]]  
[[1 2 3]]  
(a + b).shape: (4, 3),  
a + b =  
[[2 3 4]  
 [3 4 5]  
 [4 5 6]  
 [5 6 7]]
```

This is the scenario in the dense layer you built above. Adding a 1-D vector b to a (m,j) matrix.

missing

Matrix + 1-D Vector

► [Please click here if you want to experiment with any of the non-graded code.](#)